


# Introduction to Ray & Anyscale

---



# Agenda Today

- Introductions
- Overview
  - Intro your instructor
  - Why are we here?
  - What is Ray, what is Anyscale? Why Ray for ML?
- Hands On – How to Ray & Anyscale?



# Instructors & TAs

- Instructor: Charles Greer
  - Solutions Architect @ Anyscale
  - Technical Point of Contact for Koch (find me on teams!)
- TAs
  - +1 Solutions Architect: Bill Wang
  - Product Team @ Anyscale (find us on teams!)





# Why are we here today?

# Training Schedule

- **Today:** Anyscale + Ray Overview
- **In ~ + 2 Weeks: Deep Dive on Anyscale + Ray for ML dev**
  - Meeting with Maxim + Badrul next week, please send any questions / topics you'd like us to cover!
- **In ~ + 3 Weeks: Deep Dive on Anyscale + Ray for production**
  - Meeting with Dermot + Devin in the 2 weeks, please send any questions / topics you'd like us to cover!
- **Follow ups as needed (e.g., RLlib)**

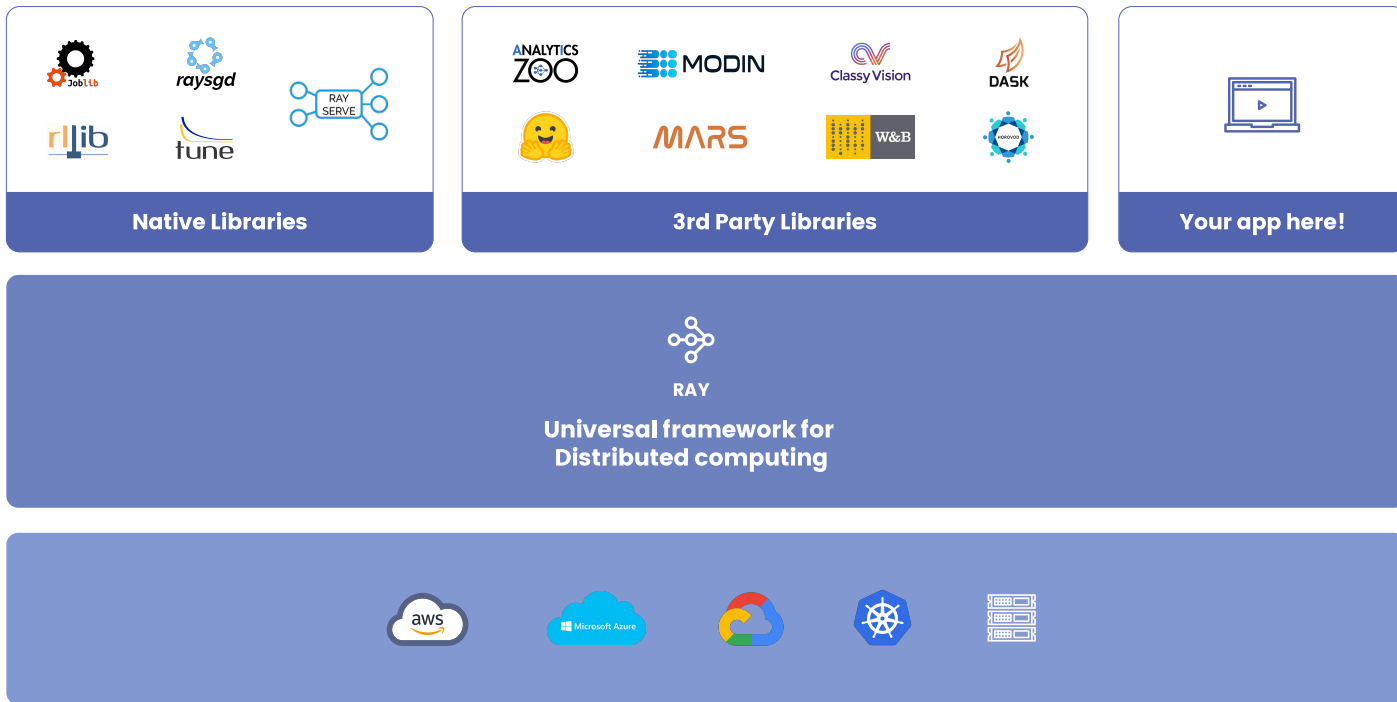




# What is Ray? What is Anyscale?

# What is Ray?

## RAY Ecosystem



# What is Ray?

- An ecosystem of tools to scale any kind of workload
  - **Ray Core:** A universal framework for executing and managing distributed computation.
  - **Ray Libraries:** A set of high-level ML-oriented libraries that leverage this united computing environment.
- A community of thousands of developers, data scientists, and ML engineers.





# Ray Core

- Distribute functions and classes across a cluster of machines without having to think about the infrastructure
- Low level tool to build apps from scratch
- Analogy: The Python Language
  - *But purpose made for scaling applications*



# Ray (Native) Libraries

- **Ray Tune** – scaling HPO with cutting edge algos
- **Ray Serve** – scaling model serving
- **Ray SGD** – scaling distributed training (multi-node + multi GPU)
- **Ray RLlib** – scaling reinforcement learning
- **Ray Data** – scaling data processing
- Analogy: The Python Standard Library
  - *But purpose made for scaling and deploying ML applications*

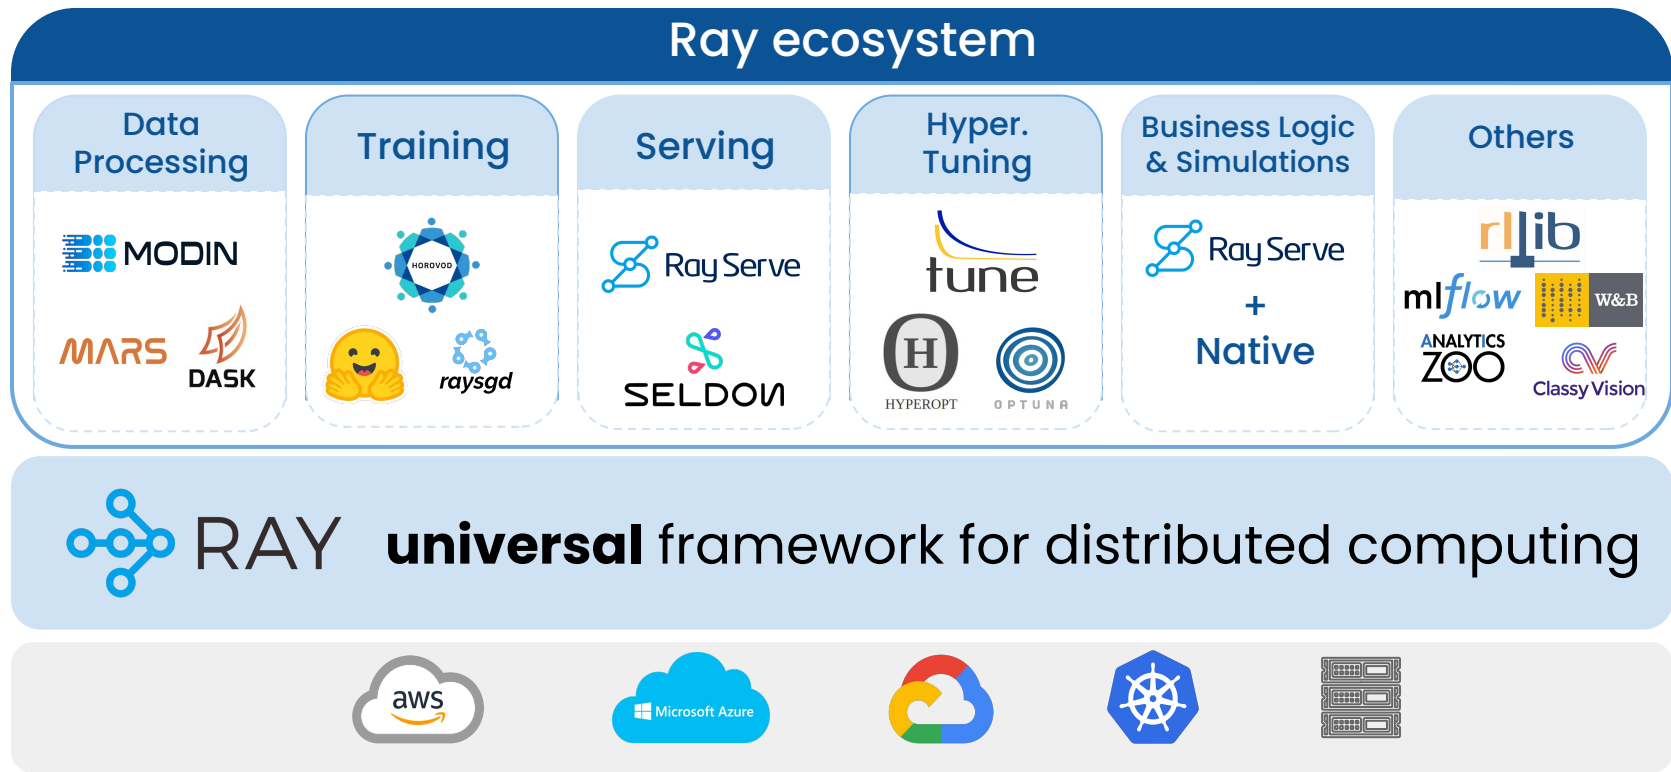


# Ray (Integrated) Libraries

- **Ray SKLearn** – scaling scikit learn
- **Ray XGBoost** – scaling XGboost
- **Ray MLFlow** – track Ray HPO / Models with MLFlow
- **Dask on Ray** – scale Dask using Ray
- Huggingface, model, classyvision, weights and biases...
- Analogy: The Python's Pypi (e.g., pip install)  
Repositories



# Ray's Ecosystem of ML and Data Libraries



## Key Point...

# Use only the parts of Ray that you need!

## Ray (Native) Libraries

- **Ray Tune** - scaling HPO with cu
- **Ray Serve** - scaling

## Ray Core

- Distribute functions and classes across a cluster of machines without having to think about the infrastructure
- Low level tool to build apps from scratch
- Analogy: The Python Language
  - But purpose made for scaling applications

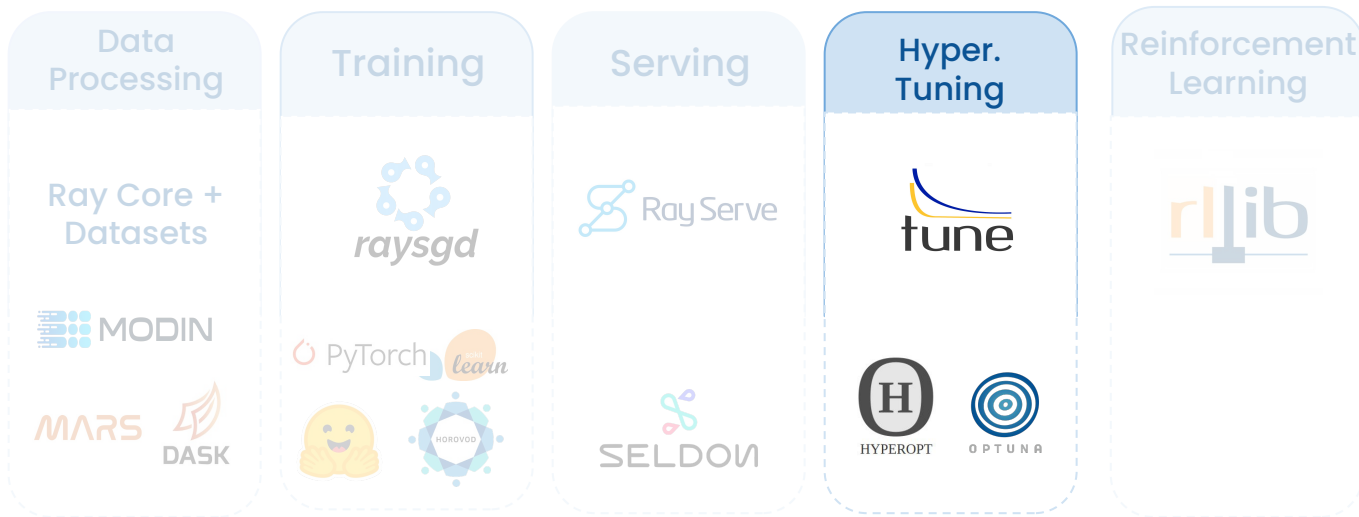
edge algos

## Ray (Integrated) Libraries

- **Ray SKLearn** - scaling scikit learn
- **Ray XGBoost** - scaling XGboost
- **Ray MLFlow** - track Ray HPO / Models with MLFlow
- **Dask on Ray** - scale Dask using Ray
- Huggingface, model, classyvision, weights and biases...
- Analogy: The Python's Pypi (e.g., pip install) Repositories



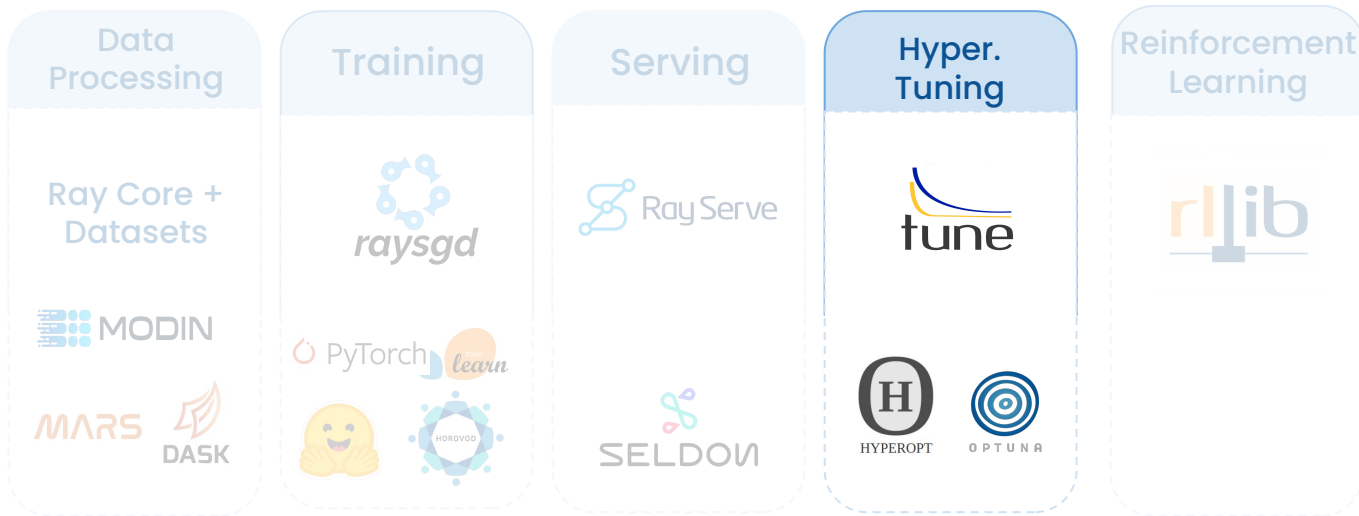
# Rich ecosystem for scaling ML workloads



Challenges in scaling hyperparameter tuning?



# Rich ecosystem for scaling ML workloads



**Integrate Ray Tune!**  
**No need to adopt entire Ray framework.**

# What is Anyscale?

## **Company founded by the Creators of Ray**

Ion Stoica (Leader RISELab, AMPLab,  
Co-Founder Databricks, Conviva),  
Michael Jordan (Leader RISELab, AMPLab),  
Robert Nishihara (Ph.D. Berkeley, Ray  
co-creator),  
Philipp Moritz (Ph.D. Berkeley, Ray co-creator)

## **Experts in Distributed Computing + ML**

ex-Uber, Stripe, Databricks, Google Brain

## **Backed by Top Tier VCs:**

A16z, NEA, Intel Capital



anyscale



# Anyscale's Mission

- Enable data teams to build ML applications faster and deploy applications more reliably by providing a managed Ray service.


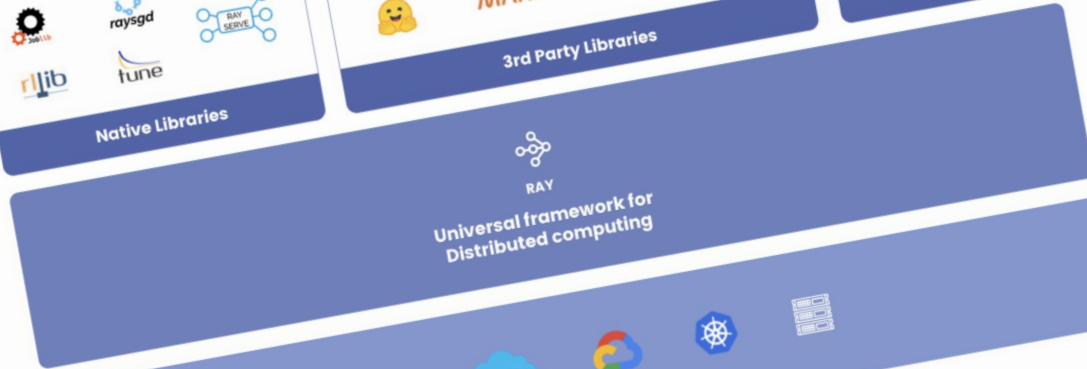


How do we do this?

- Enable Koch to focus on ML applications that make a difference for the business, not infrastructure to get a service out the door.
- Provide simple ways of tapping into scaling (e.g., remote connect to a Ray cluster)
- Provide robust APIs for CI/CD + deployment



## What is Ray?

### RAY Ecosystem



## RAY Ecosystem



Native Libraries



3rd Party Libraries



Your app here!



RAY

Universal framework for  
Distributed computing



## RAY Ecosystem



Native Libraries



3rd Party Libraries



Your app here!



RAY

Universal framework for  
Distributed computing

Anyscale leverages the same Ray  
Open Source APIs

### **Workflow Advantages:**

- Simple access to scale in dev
- Streamlined APIs to move to production

### **System Advantages:**

- Robust APIs for managing 10s - 100s of clusters
- Proprietary Cluster Manager
- Optimized Ray Runtime
- Integrated metrics & monitoring





# Hands-on Demo & Lab Using Ray on Anyscale

# What is Ray – 3 Key Ideas

Execute **functions** remotely as **tasks**, and instantiate **classes** remotely as **actors**

- **Support both stateful and stateless computations**

Asynchronous execution using futures

- **Enable parallelism**

Distributed (immutable) object store

- **Efficient communication** (send arguments by reference)



# What is Ray? API looks like:



Function → Task

Class → Actor

Object → (Distributed) Object



# What is Ray – API

```
def read_array(file):  
    # read array a from file  
    return a
```

```
def add(a, b):  
    return np.add(a, b)
```

```
a = read_array(file1)  
b = read_array(file2)  
sum = add(a, b)
```

```
class Counter(object):  
    def __init__(self):  
        self.value = 0  
    def inc(self):  
        self.value += 1  
    return self.value
```

```
c = Counter()  
c.inc()  
c.inc()
```





# What is Ray - API

```
@ray.remote
```

```
def read_array(file):  
    # read array a from file  
    return a
```

```
@ray.remote
```

```
def add(a, b):  
    return np.add(a, b)
```

```
a = read_array(file1)  
b = read_array(file2)  
sum = add(a, b)
```

```
@ray.remote
```

```
class Counter(object):  
    def __init__(self):  
        self.value = 0  
    def inc(self):  
        self.value += 1  
    return self.value
```

```
c = Counter()  
c.inc()  
c.inc()
```



# What is Ray - API

```
@ray.remote
def read_array(file):
    # read array a from file
    return a
```

```
@ray.remote
def add(a, b):
    return np.add(a, b)
```

```
ref1 = read_array.remote(file1)
ref2 = read_array.remote(file2)
ref = add.remote(ref1, ref2)
sum = ray.get(ref)
```

```
@ray.remote(num_gpus=1)
class Counter(object):
    def __init__(self):
        self.value = 0
    def inc(self):
        self.value += 1
        return self.value
```

```
c = Counter.remote()
ref4 = c.inc.remote()
ref5 = c.inc.remote()
```



# What is Ray – API (Actor Handles)

*Invoke actor methods from other tasks/actors/applications.*

```
@ray.remote(num_gpus=1)
class Counter(object):
    def __init__(self):
        self.value = 0
    def inc(self):
        self.value += 1
        return self.value
```

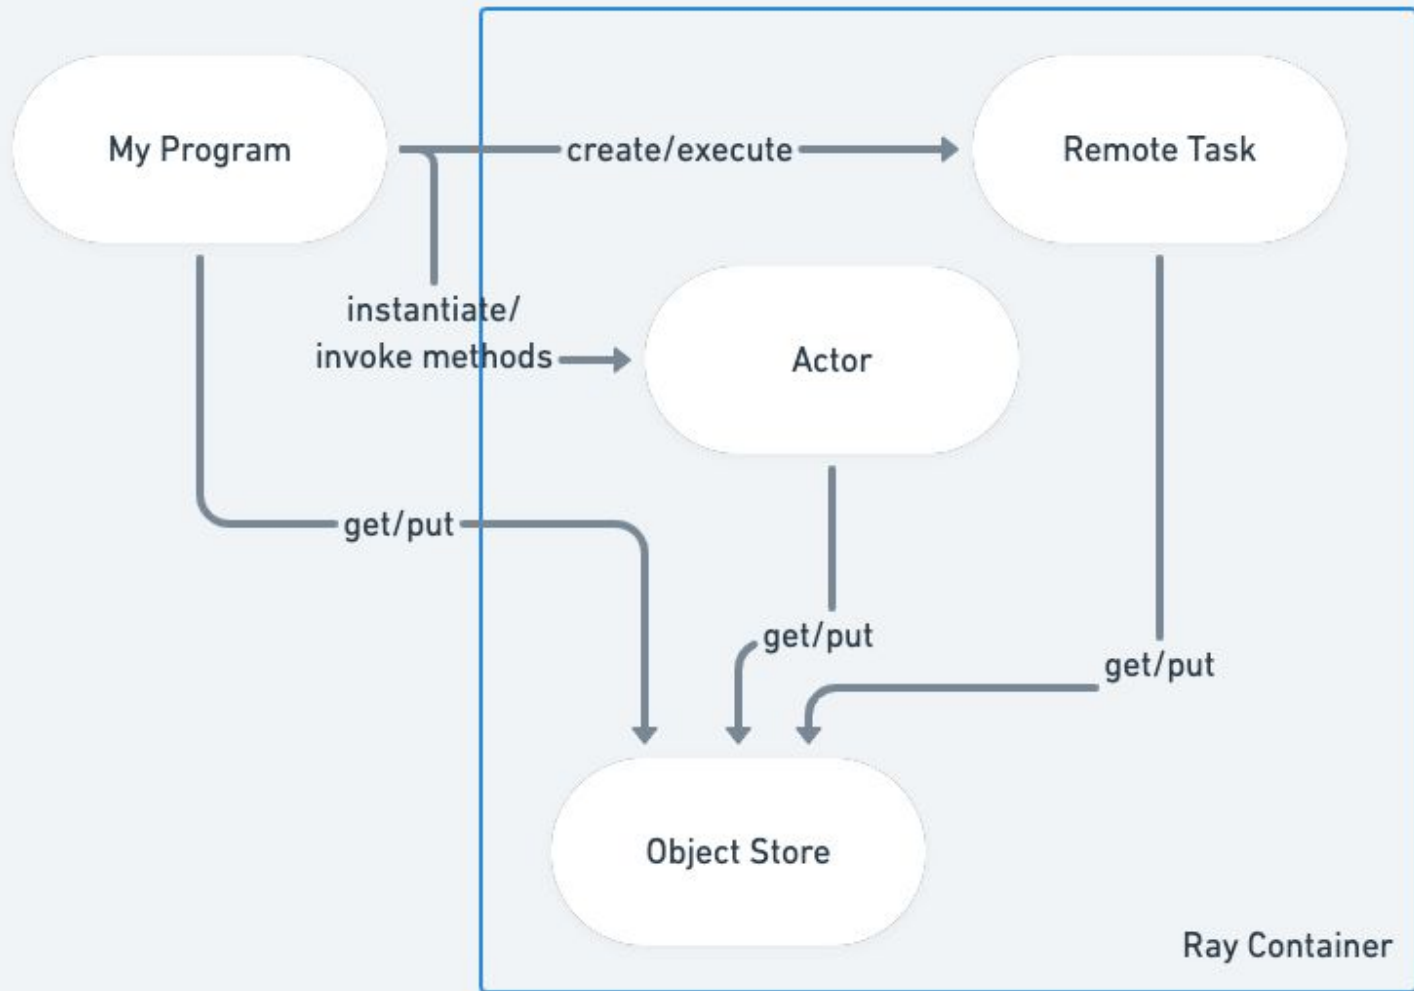
```
c = Counter.remote()
id = c.inc.remote()
```

```
# Use the actor from a
# different task
```

```
@ray.remote
def use_actor(c):
    id = c.inc.remote()
    ray.get(id)
```

```
use_actor.remote(c)
```





# Ray Application Lifecycle


- Experiment/Design on a Single Node
  - Jupyter Notebook on Anyscale
  - Laptop and local ray
- Optimize and Scale
- Automate and Deploy



# Let us Code

- Tasks
- Actors
- Remote calls
- Put and get
- Best Practices
- Scaling






# What is Ray? Hello World...

```
import ray
ray.init()

@ray.remote
def f(x):
    return x * x
```

```
futures = [f.remote(i) for i in range(4)]
print(ray.get(futures)) # [0, 1, 4, 9]
```





# What is Ray? Hello World...

```
@ray.remote(num_cpus=1)
```

```
class Worker:
```

```
    def work(self):
```

```
        return "done"
```

```
@ray.remote(num_cpus=1)
```

```
class Supervisor:
```

```
    def __init__(self):
```

```
        self.workers = [Worker.remote() for _ in range(3)]
```


```
    def work(self):
```

```
        return ray.get([w.work.remote() for w in self.workers])
```

```
ray.init()
```

```
sup = Supervisor.remote()
```

```
print(ray.get(sup.work.remote())) # outputs ['done', 'done', 'done']
```



# Additional Resources Ray

## Documentation ([docs.ray.io](https://docs.ray.io))

*Quick start example, reference guides, etc*

## Forums ([discuss.ray.io](https://discuss.ray.io))

*Learn / share with broader Ray community, including core team*

## Ray Slack

*Connect with the Ray team and community*

## Anyscale Support

*Find us on Teams, on the Open Source Ray Slack, or over email.*





**Charles Greer**  
**Bill Wang**  
**Bill Chambers**  
**Javier Redondo**  
**Tricia Fu**  
**Will Drevo**

---