

Anyscale Devops



Agenda Today

- Anyscale Systems Overview
- CLI – Anyscale from your laptop
- SDK – Scripting Anyscale
- Application Lifecycle
- Ray Serve



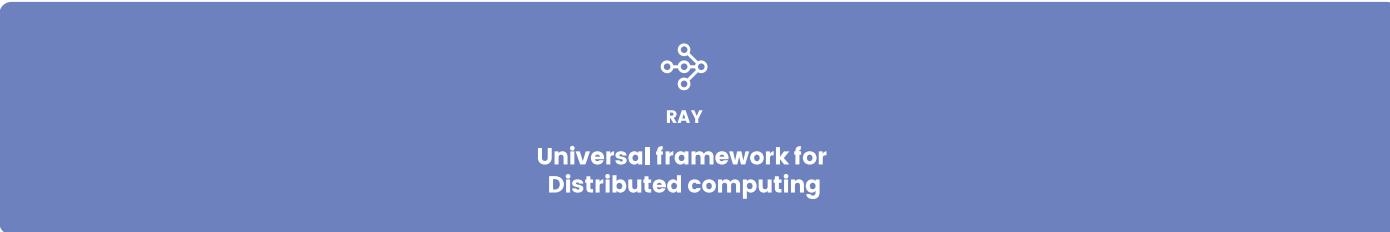
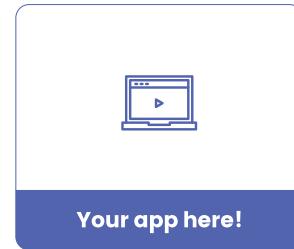
Instructors & TAs

- Instructor: Charles Greer
 - Solutions Architect @ Anyscale
 - Technical Point of Contact for Koch (find me on teams!)
- TAs
 - Product Team @ Anyscale



What is Ray?

RAY Ecosystem



What is Ray?

- A place to run tasks in a massively parallel way.
- A framework for executing and managing distributed computation.
- A set of high-level ML-oriented libraries that leverage this united computing environment.
- A community of thousands of developers and ML engineers.



What is Anyscale?

Company founded by the Creators of Ray

Ion Stoica (Leader RISELab, AMPLab,
Co-Founder Databricks, Conviva),
Michael Jordan (Leader RISELab, AMPLab),
Robert Nishihara (Ph.D. Berkeley, Ray
co-creator),
Philipp Moritz (Ph.D. Berkeley, Ray co-creator)

Experts in Distributed Computing + ML

ex-Uber, Stripe, Databricks, Google Brain

Backed by Top Tier VCs:

A16z, NEA, Intel Capital



anyscale



Why are you here?

Time to get Real

- Over the past 3 months we've been working with the Koch team on a couple of different use cases.
- Now we see how Anyscale ML workflows are incorporated into the Enterprise

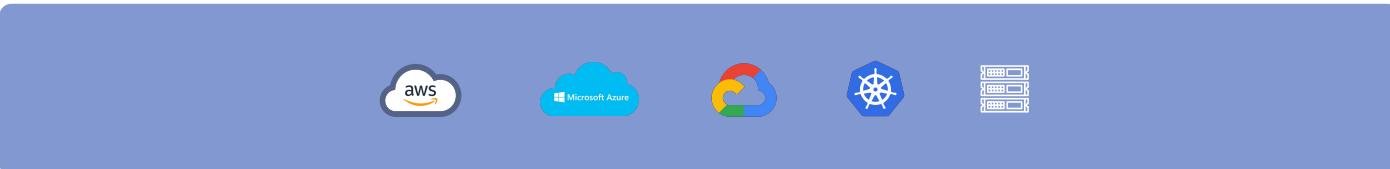
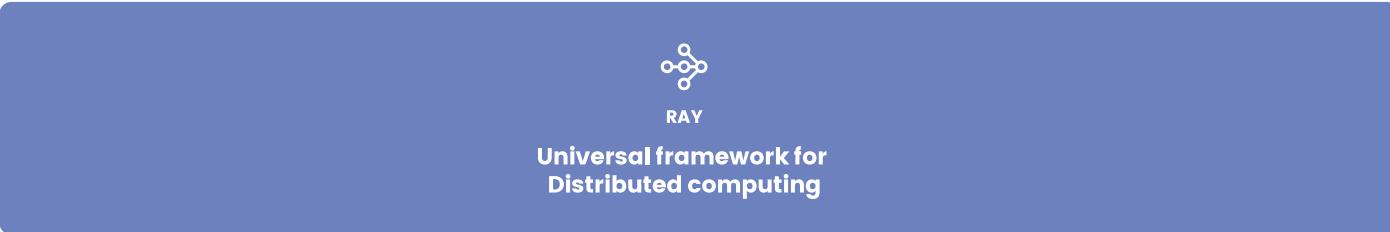
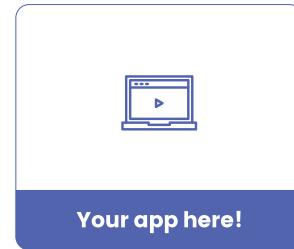




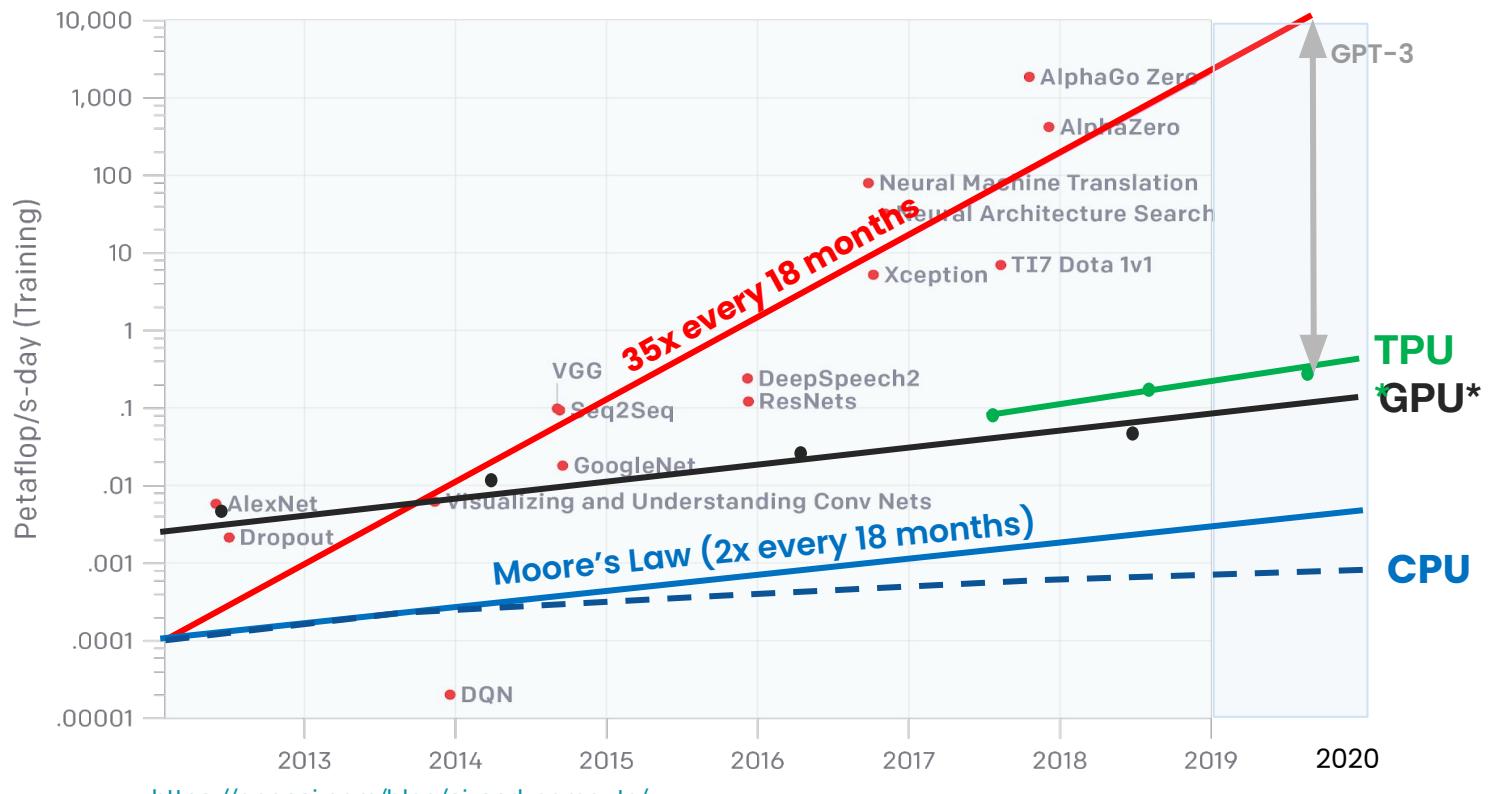
**Anyscale's goal:
Koch focuses on building
ML applications, not
managing infrastructure**

What is Ray?

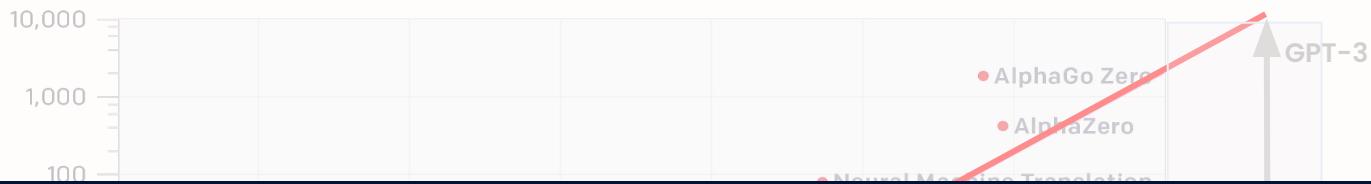
RAY Ecosystem



Specialized hardware is also not enough



Specialized hardware is also not enough



No way out but to distribute!



Stitching together different frameworks to go end-to-end?

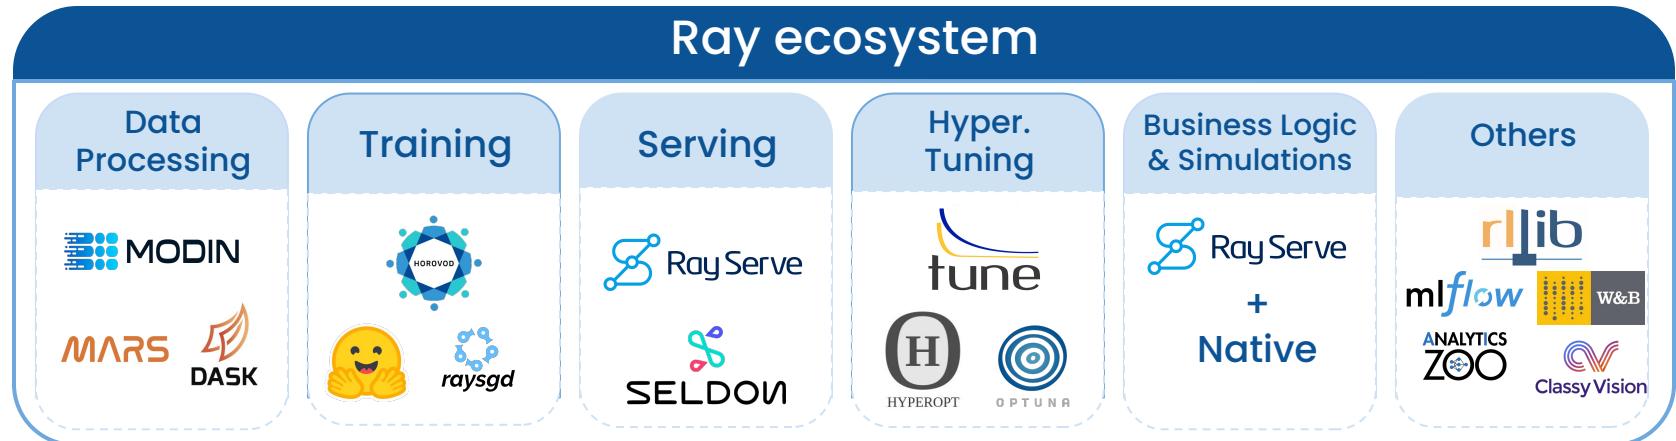
Ease of development



Generality



Ray's Ecosystem of ML and Data Libraries



RAY universal framework for distributed computing





Anyscale's Mission

- 01 Enable users to focus on apps instead of infrastructure.
- 02 Provide rapid iteration + reduce time to market for new products
- 03 Support variable infrastructure and clouds
 - make compute magical

What is Ray - 3 Key Ideas

Execute **functions** remotely as **tasks**, and
instantiate **classes** remotely as **actors**

- **Support both stateful and stateless computations**

Asynchronous execution using futures

- **Enable parallelism**

Distributed (immutable) object store

- **Efficient communication** (send arguments by reference)



What is Ray - API

```
def read_array(file):
    # read array a from file
    return a

def add(a, b):
    return np.add(a, b)

a = read_array(file1)
b = read_array(file2)
sum = add(a, b)
```

```
class Counter(object):
    def __init__(self):
        self.value = 0
    def inc(self):
        self.value += 1
    return self.value

c = Counter()
c.inc()
c.inc()
```



What is Ray - API

```
@ray.remote  
def read_array(file):  
    # read array a from file  
    return a  
  
@ray.remote  
def add(a, b):  
    return np.add(a, b)  
  
a = read_array(file1)  
b = read_array(file2)  
sum = add(a, b)
```

```
@ray.remote  
class Counter(object):  
    def __init__(self):  
        self.value = 0  
    def inc(self):  
        self.value += 1  
    return self.value  
  
c = Counter()  
c.inc()  
c.inc()
```



What is Ray - API

```
@ray.remote  
def read_array(file):  
    # read array a from file  
    return a  
  
@ray.remote  
def add(a, b):  
    return np.add(a, b)  
  
ref1 = read_array.remote(file1)  
ref2 = read_array.remote(file2)  
ref = add.remote(ref1, ref2)  
sum = ray.get(ref)
```

```
@ray.remote(num_gpus=1)  
class Counter(object):  
    def __init__(self):  
        self.value = 0  
    def inc(self):  
        self.value += 1  
        return self.value  
  
c = Counter.remote()  
ref4 = c.inc.remote()  
ref5 = c.inc.remote()
```



What is Ray - API (Actor Handles)

Invoke actor methods from other tasks/actors/applications.

```
@ray.remote(num_gpus=1)
class Counter(object):
    def __init__(self):
        self.value = 0
    def inc(self):
        self.value += 1
        return self.value

c = Counter.remote()
id = c.inc.remote()
```

```
# Use the actor from a
# different task

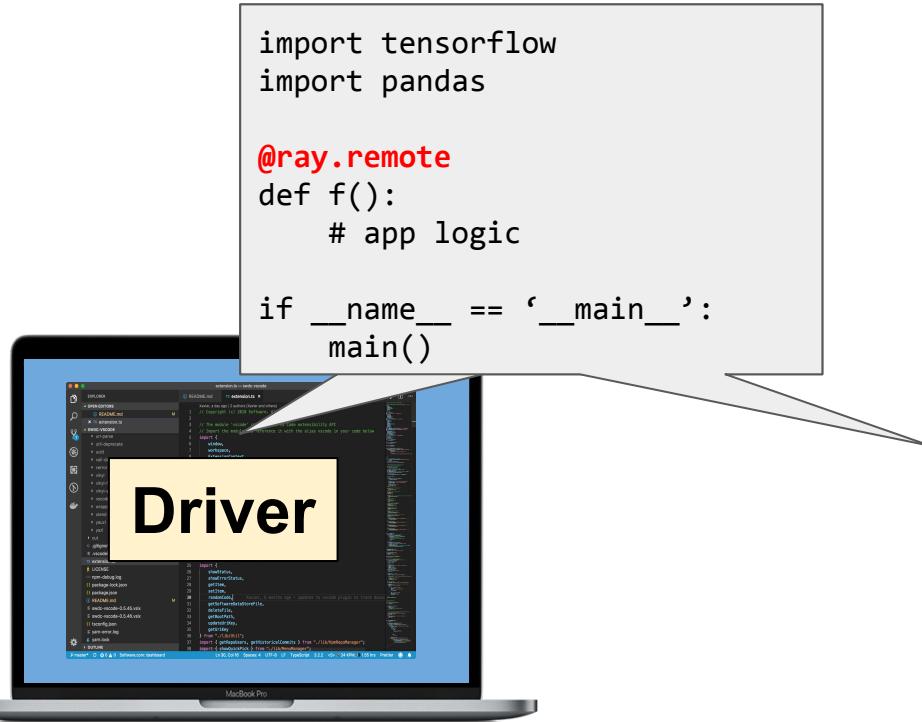
@ray.remote
def use_actor(c):
    id = c.inc.remote()
    ray.get(id)

use_actor.remote(c)
```



Streamlined Workflows

```
ray.init()
```



Ray cluster



Streamlined Workflows

Environments

```
ray.init(runtime_env=
{
    'pip': ['tensorflow', 'pandas'],
})
)
```



Streamlined Workflows

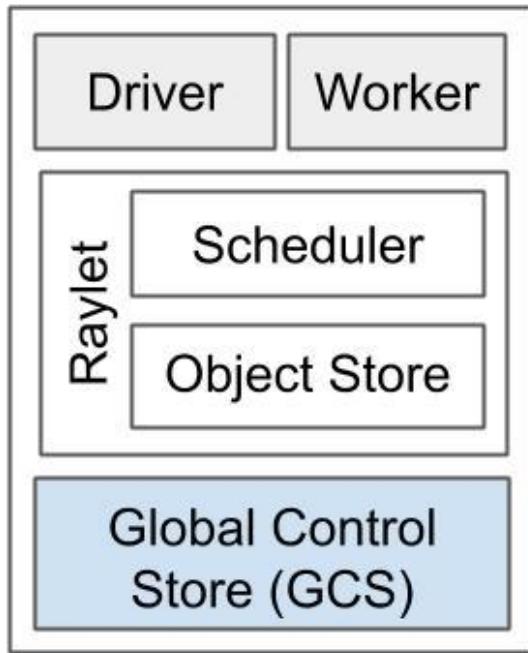
Environments

```
ray.init(runtime_env=
{
    'pip': ['tensorflow', 'pandas'],
    'working_dir': '/home/my-project'
})
)
```

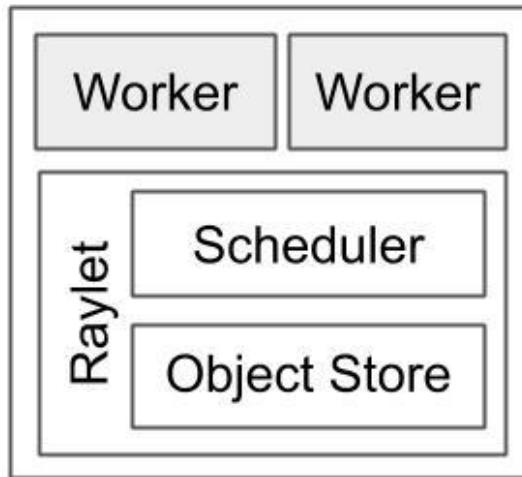


Ray Architecture

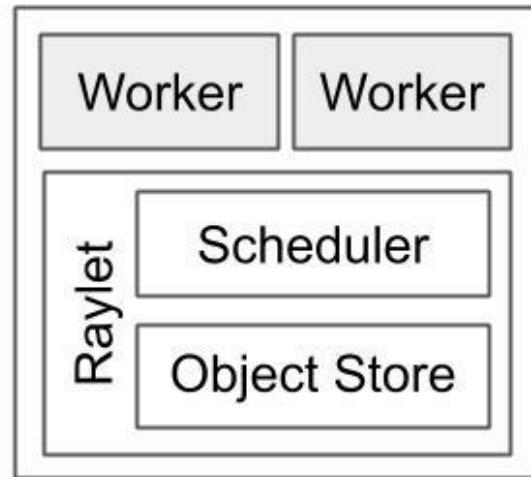
Head node



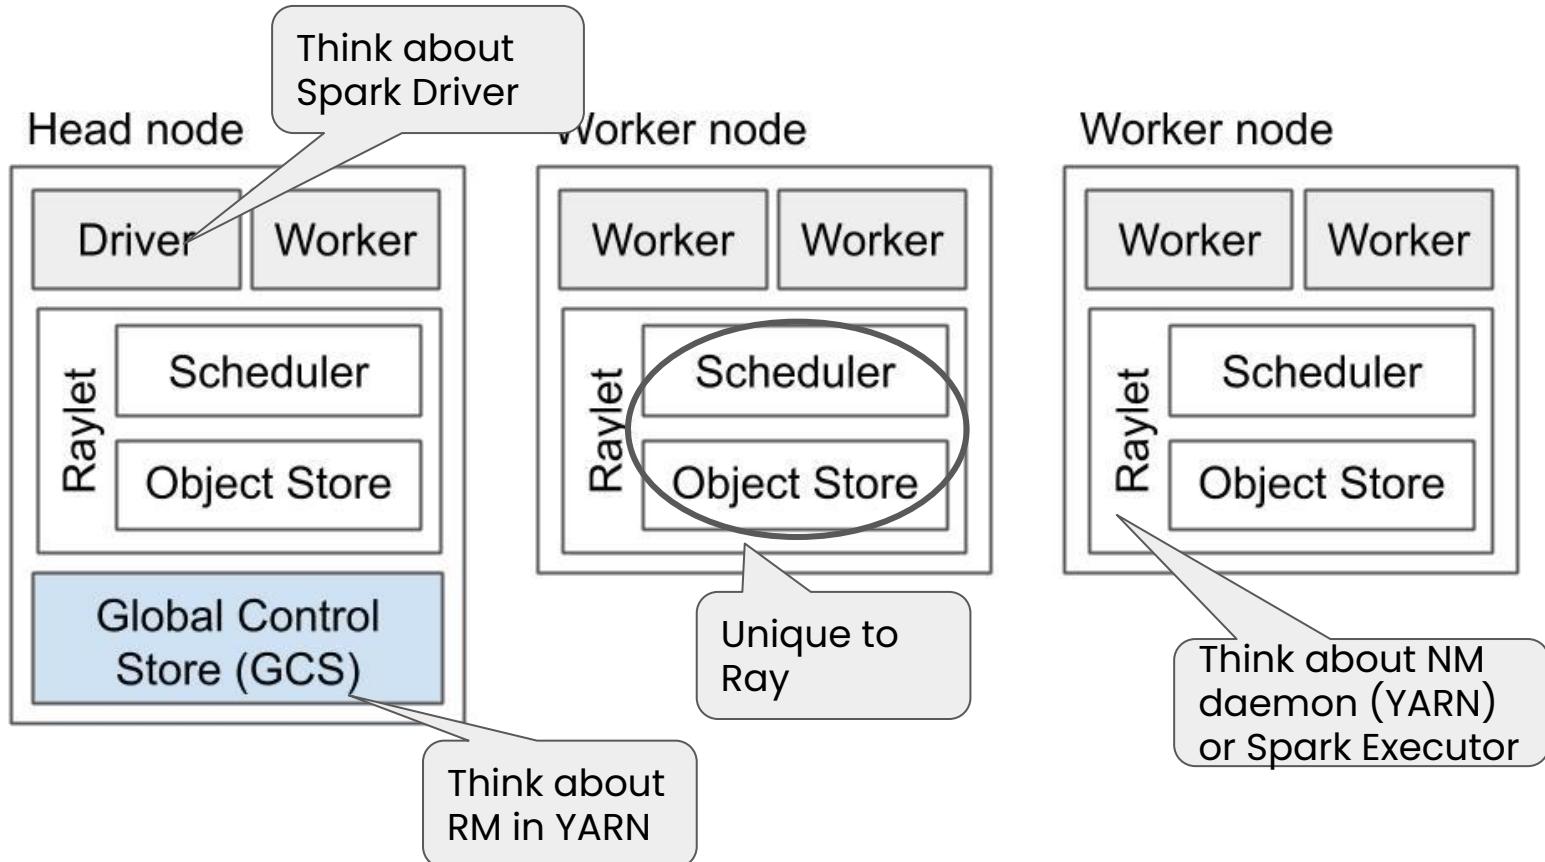
Worker node



Worker node



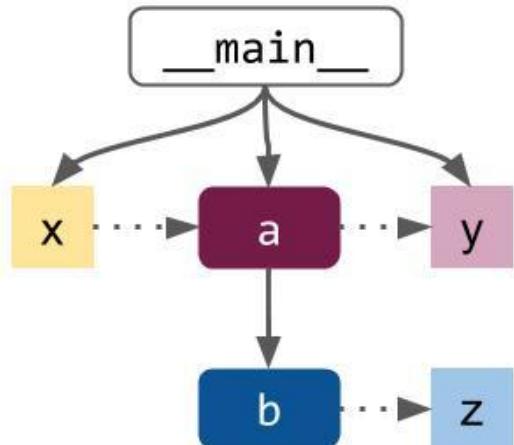
What is Ray? A Cluster Looks Like...



What is Ray? API -> Cluster

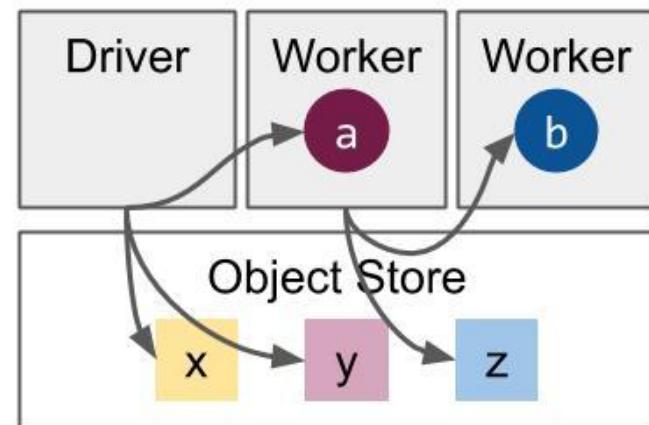
```
@ray.remote  
def b():  
    return  
  
@ray.remote  
def a(dep):  
    z = b.remote()  
  
x = ray.put(...)  
y = a.remote(x)
```

Program



Task graph

→ Ownership
→ Dependency



Physical execution



Ray Application Lifecycle

- Experiment/Design on a Single Node
 - Jupyter Notebook on Anyscale
 - Laptop and local ray
- Optimize and Scale
- **Automate and Deploy**

github: anyscale/training-one



CLI and Anyscale Resources

- pip install anyscale
- anyscale init
- The project - an ambiguous term
 - A Project in Anyscale
 - Working Directory
- Python environment



Anyscale CLI - running commands

anyscale run

- Still experimental, let's return to this later.



Anyscale SDK

- Building Cluster Environments
- Launching Clusters
- Parallelizing SDK tasks (with ray!)



Application Lifecycle

- Project Structure
- Entry Points
 - Unit testing
 - Integration Testing
 - Application Interface



Ray Serve

- FastAPI
- Replication
- Rate Testing with Locust



Additional Resources

Documentation ([docs.ray.io](#))

Quick start example, reference guides, etc

Forums ([discuss.ray.io](#))

Learn / share with broader Ray community, including core team

Ray Slack

Connect with the Ray team and community

